

CS 2601 Linear and Convex Optimization

6. Gradient descent (part 1)

Bo Jiang

John Hopcroft Center for Computer Science
Shanghai Jiao Tong University

Fall 2024

Unconstrained optimization problems

Consider an unconstrained, smooth convex optimization problem

$$\min_{\mathbf{x}} f(\mathbf{x})$$

where f is convex and differentiable on $\mathbb{R}^n$.

The optimal solution satisfies the first-order optimality condition

$$\nabla f(\mathbf{x}^*) = \mathbf{0}$$

In some rare cases, this yields closed-form solutions, e.g.

$$\min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$



has closed-form solution

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

But in most cases we need numerical algorithms.

Descent method

- 1: choose initial point $\mathbf{x}_0 \in \mathbb{R}^n$
- 2: **repeat**
- 3: choose **descent direction** $\mathbf{d}_k \in \mathbb{R}^n$ and **step size** $t_k > 0$
- 4: $\mathbf{x}_{k+1} = \mathbf{x}_k + t_k \mathbf{d}_k$ s.t. $f(\mathbf{x}_{k+1}) < f(\mathbf{x}_k)$
- 5: **until** stopping criterion is satisfied



Questions

- How to choose $\mathbf{d}_k$ and t_k ?
- Does $\mathbf{x}_k$ converge to $\mathbf{x}^*$?

Descent direction

$\mathbf{d}_k$ is a **descent direction** at $\mathbf{x}_k$ if for all small enough $t > 0$

$$g(t) \triangleq f(\mathbf{x}_k + t\mathbf{d}_k) < f(\mathbf{x}_k) = g(0)$$

For differentiable f (not necessarily convex),

- if $\mathbf{d}_k$ is a descent direction, then $g'(0) = \mathbf{d}_k^T \nabla f(\mathbf{x}_k) \leq 0$;
- if $g'(0) = \mathbf{d}_k^T \nabla f(\mathbf{x}_k) < 0$, then $\mathbf{d}_k$ is a descent direction.

For convex f , by the first-order condition for convexity,

$$f(\mathbf{x}_k) > f(\mathbf{x}_k + t\mathbf{d}_k) \geq f(\mathbf{x}_k) + t\mathbf{d}_k^T \nabla f(\mathbf{x}_k).$$

$\mathbf{d}_k^T \nabla f(\mathbf{x}_k) < 0$ is also **necessary** for $\mathbf{d}_k$ to be a descent direction.


For convex differentiable f ,

$$\mathbf{d}_k \text{ is a descent direction} \iff \mathbf{d}_k^T \nabla f(\mathbf{x}_k) < 0$$


Gradient descent

Choose $\mathbf{d}_k = -\nabla f(\mathbf{x}_k)$, $\mathbf{d}_k^T \nabla f(\mathbf{x}_k) = -\|\nabla f(\mathbf{x}_k)\|_2^2 < 0$ unless $\nabla f(\mathbf{x}_k) = \mathbf{0}$.

$$\mathbf{x}_{k+1} = \mathbf{x}_k - t_k \nabla f(\mathbf{x}_k)$$



level curves of $f(x_1, x_2) = \frac{x_1^2}{4} + x_2^2$



Question. What happens if $\nabla f(\mathbf{x}_k) = \mathbf{0}$?


Max-rate descending direction

$-\nabla f(\mathbf{x}_k)$ is the direction of fastest rate of decrease of f at $\mathbf{x}_k$


- If $\|\mathbf{d}_k\|_2 = 1$,

$$\lim_{t \downarrow 0} \frac{f(\mathbf{x}_k) - f(\mathbf{x}_k + t\mathbf{d}_k)}{t} = -\mathbf{d}_k^T \nabla f(\mathbf{x}_k) \leq \|\nabla f(\mathbf{x}_k)\|_2$$

with equality iff $\mathbf{d}_k = -\nabla f(\mathbf{x}_k) / \|\nabla f(\mathbf{x}_k)\|_2$



level curves of $f(x_1, x_2) = \frac{x_1^2}{4} + x_2^2$



Gradient descent algorithm

- 1: initialization $\mathbf{x} \leftarrow \mathbf{x}_0 \in \mathbb{R}^n$
- 2: **while** $\|\nabla f(\mathbf{x})\| > \delta$ **do**
- 3: $\mathbf{x} \leftarrow \mathbf{x} - t \nabla f(\mathbf{x})$
- 4: **end while**
- 5: **return** $\mathbf{x}$

Step size (aka **learning rate** in machine learning)

- the above algorithm uses constant step size t for all iterations
- there are other methods for choosing t for each iteration, e.g. **exact line search**, **backtracking line search**







Stopping criterion

- ideally, stop if $\nabla f(\mathbf{x}) = 0$ (optimality condition), but impractical
- more practical: stop when $\|\nabla f(\mathbf{x})\| \leq \delta$ for some small δ
- other criteria: $|f(\mathbf{x}_{\text{new}}) - f(\mathbf{x}_{\text{old}})| \leq \delta$, $\frac{|f(\mathbf{x}_{\text{new}}) - f(\mathbf{x}_{\text{old}})|}{|f(\mathbf{x}_{\text{old}})|} \leq \delta$, ...
- in practice, also stop if maximum # of iterations is reached

Large vs. small step size

Consider constant step size. How large should the step size be?

- Too large: may oscillate and diverge
- Too small: may be too slow
- “Just right”: fast convergence



1D example

Consider $f(x) = \frac{1}{2}ax^2$, where $a > 0$.

- gradient step

$$x_{k+1} = x_k - tf'(x_k) = (1 - at)x_k$$

- descent condition


$$f(x_{k+1}) < f(x_k) \iff |1 - at| < 1 \iff 0 < t < \frac{2}{a}$$

- $x_k = (1 - at)^k x_0 \rightarrow x^* = 0$ geometrically for such t

Note f satisfies

- $|f'(x) - f'(y)| = a|x - y|$
- $f''(x) = a$

f' is so-called **Lipschitz continuous**
and t is roughly the order of $\frac{1}{a}$.



Lipschitz continuity

A function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is **Lipschitz continuous** with **Lipschitz constant** $L > 0$, or simply **L -Lipschitz**, if

$$\|f(\mathbf{x}) - f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|, \quad \forall \mathbf{x}, \mathbf{y}$$

Note. Lipschitz continuity can be defined with respect to any norms. But we will assume the norms in the above definition are the 2-norms in $\mathbb{R}^n$ and $\mathbb{R}^m$, respectively, unless stated otherwise.

Note. Lipschitz continuity implies uniform continuity.

Example. $f(x) = ax$ is $|a|$ -Lipschitz, $|f(x) - f(y)| = |a| \cdot |x - y|$

Example. $f(x) = |x|$ is 1-Lipschitz, $|f(x) - f(y)| = ||x| - |y|| \leq |x - y|$

Example. $f(\mathbf{x}) = \mathbf{a}^T \mathbf{x}$ is $\|\mathbf{a}\|$ -Lipschitz, $|\mathbf{a}^T \mathbf{x} - \mathbf{a}^T \mathbf{y}| \leq \|\mathbf{a}\| \cdot \|\mathbf{x} - \mathbf{y}\|$ by the Cauchy-Schwarz inequality.

Lipschitz continuity (cont'd)

Example. Let $\mathbf{Q} = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$. $f(\mathbf{x}) = \mathbf{Q}\mathbf{x} = (x_1, 2x_2)$ is 2-Lipschitz.

$$f(\mathbf{x}) - f(\mathbf{y}) = (x_1 - y_1, 2x_2 - 2y_2) = (d_1, 2d_2)$$

$$\|f(\mathbf{x}) - f(\mathbf{y})\| = \sqrt{d_1^2 + 4d_2^2} \leq 2\sqrt{d_1^2 + d_2^2} = 2\|\mathbf{x} - \mathbf{y}\|$$

More generally, $f(\mathbf{x}) = \mathbf{Q}\mathbf{x}$ with $\mathbf{Q} \succeq \mathbf{O}$ is $\lambda_{\max}(\mathbf{Q})$ -Lipschitz, where $\lambda_{\max}(\mathbf{Q})$ is the largest eigenvalue of $\mathbf{Q}$ ¹.

Proof. Let $\mathbf{d} = \mathbf{x} - \mathbf{y}$. By slide 32 of §2,

$$\|f(\mathbf{x}) - f(\mathbf{y})\| = \|\mathbf{Q}\mathbf{d}\| = \sqrt{\mathbf{d}^T \mathbf{Q}^2 \mathbf{d}} \leq \sqrt{\lambda_{\max}(\mathbf{Q}^2) \|\mathbf{d}\|^2} = \lambda_{\max}(\mathbf{Q}) \|\mathbf{x} - \mathbf{y}\|$$

The last equality uses the fact $\lambda_{\max}(\mathbf{Q}^2) = \lambda_{\max}^2(\mathbf{Q})$.

¹For general $\mathbf{Q}$, $f(\mathbf{x}) = \mathbf{Q}\mathbf{x}$ is $\sigma_{\max}(\mathbf{Q})$ -Lipschitz, where $\sigma_{\max}(\mathbf{Q}) = \sqrt{\lambda_{\max}(\mathbf{Q}^T \mathbf{Q})}$ is the largest singular value of $\mathbf{Q}$.

L -smoothness

A function is L -smooth if it is differentiable and its gradient is L -Lipschitz, i.e.

$$\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|, \quad \forall \mathbf{x}, \mathbf{y}$$

Note. L upper bounds the rate of change of ∇f

Example. $f(x) = \frac{1}{2}ax^2$ is $|a|$ -smooth, since $f'(x) = ax$ is $|a|$ -Lipschitz

Example. $f(\mathbf{x}) = \frac{1}{2}\mathbf{x}^T \mathbf{Q} \mathbf{x}$ with $\mathbf{Q} \succeq \mathbf{O}$ is $\lambda_{\max}(\mathbf{Q})$ -smooth, since $\nabla f(\mathbf{x}) = \mathbf{Q} \mathbf{x}$ is $\lambda_{\max}(\mathbf{Q})$ -Lipschitz.

With $\mathbf{Q} = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$, we obtain $f(\mathbf{x}) = \frac{1}{2}x_1^2 + x_2^2$ is 2-smooth.

Lemma. A twice continuously differentiable **convex** $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is L -smooth iff $\nabla^2 f(\mathbf{x}) \preceq L\mathbf{I}$, meaning $L\mathbf{I} - \nabla^2 f(\mathbf{x}) \succeq \mathbf{O}$, or equivalently $\lambda_{\max}(\nabla^2 f(\mathbf{x})) \leq L$.

Appendix: Second-order condition for L -smoothness

Lemma. A twice continuously differentiable $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is L -smooth iff for any $\mathbf{x}$, $-L\mathbf{I} \preceq \nabla^2 f(\mathbf{x}) \preceq L\mathbf{I}$, or equivalently $|\lambda| \leq L$ for all eigenvalues λ of $\nabla^2 f(\mathbf{x})$.

Proof. “ $\Leftarrow$ ”. Assume $-L\mathbf{I} \preceq \nabla^2 f(\mathbf{x}) \preceq L\mathbf{I}$ for all $\mathbf{x}$. Let $g(\mathbf{x}) = \mathbf{u}^T \nabla f(\mathbf{x})$. Then $\nabla g(\mathbf{x}) = \nabla^2 f(\mathbf{x})\mathbf{u}$.

- By slide 32 of §2, $\|\nabla g(\mathbf{x})\| \leq \lambda_{\max}(\nabla^2 f(\mathbf{x}))\|\mathbf{u}\| \leq L\|\mathbf{u}\|$.
- By the Mean Value Theorem,

$$\mathbf{u}^T [\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})] = g(\mathbf{x}) - g(\mathbf{y}) = \nabla g(\mathbf{z})^T (\mathbf{x} - \mathbf{y}) \leq L\|\mathbf{u}\| \cdot \|\mathbf{x} - \mathbf{y}\|$$

Setting $\mathbf{u} = \nabla f(\mathbf{x}) - \nabla f(\mathbf{y})$ yields $\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|$.

“ $\Rightarrow$ ”. Assume f is L -smooth. Let $\mathbf{d}$ be an eigenvector of $\nabla^2 f(\mathbf{x})$ with associated eigenvalue λ . By L -smoothness,

$$\|\nabla f(\mathbf{x} + t\mathbf{d}) - \nabla f(\mathbf{x})\| \leq L\|t\mathbf{d}\| = tL\|\mathbf{d}\|$$


Dividing both sides by t and letting $t \rightarrow 0$,

$$|\lambda| \cdot \|\mathbf{d}\| = \|\nabla^2 f(\mathbf{x})\mathbf{d}\| \leq L\|\mathbf{d}\| \implies |\lambda| \leq L$$

Quadratic upper bound

Lemma. If f is L -smooth, then

$$f(\mathbf{y}) \leq f(\mathbf{x}) + \nabla f(\mathbf{x})^T (\mathbf{y} - \mathbf{x}) + \frac{L}{2} \|\mathbf{y} - \mathbf{x}\|^2$$



Note. The upper bound does **not** assume the convexity of f .

If $\nabla^2 f(\mathbf{x}) \preceq L\mathbf{I}$, the upper bound is intuitive from the Taylor expansion

$$f(\mathbf{y}) = f(\mathbf{x}) + \nabla f(\mathbf{x})^T (\mathbf{y} - \mathbf{x}) + \frac{1}{2} (\mathbf{y} - \mathbf{x})^T \nabla^2 f(\mathbf{x} + \theta(\mathbf{y} - \mathbf{x})) (\mathbf{y} - \mathbf{x})$$

for some $\theta \in [0, 1]$. (Check $f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T \mathbf{Q} \mathbf{x}$)

Proof

First prove the 1D case. Let $g(t)$ be L_g -smooth, $|g'(t) - g'(s)| \leq L_g|t - s|$.

$$\begin{aligned} g(1) &= g(0) + \int_0^1 g'(t) dt \\ &= g(0) + g'(0) + \int_0^1 [g'(t) - g'(0)] dt \\ &\leq g(0) + g'(0) + \int_0^1 L_g t dt \quad \text{since } |g'(t) - g'(0)| \leq L_g t \\ &= g(0) + g'(0) + \frac{1}{2} L_g \end{aligned}$$

For the general case, apply the above to $g(t) = f(\mathbf{x} + t\mathbf{d})$ with $\mathbf{d} = \mathbf{y} - \mathbf{x}$ and $L_g = L\|\mathbf{d}\|^2$. To check the smoothness of g ,

$$\begin{aligned} |g'(t) - g'(s)| &= |[\nabla f(\mathbf{x} + t\mathbf{d}) - \nabla f(\mathbf{x} + s\mathbf{d})]^T \mathbf{d}| \\ &\leq \|\nabla f(\mathbf{x} + t\mathbf{d}) - \nabla f(\mathbf{x} + s\mathbf{d})\| \cdot \|\mathbf{d}\| \quad \text{Cauchy-Schwarz} \\ &\leq (t - s)L\|\mathbf{d}\|^2 \quad \text{\textit{f is L-smooth}} \end{aligned}$$

Consequence of quadratic upper bound

For L -smooth f , the sequence $\{\mathbf{x}_k\}$ produced by gradient descent satisfies

$$f(\mathbf{x}_{k+1}) \leq f(\mathbf{x}_k) - t \left(1 - \frac{Lt}{2}\right) \|\nabla f(\mathbf{x}_k)\|^2$$

Proof. Plugging in $\mathbf{x} = \mathbf{x}_k$ and $\mathbf{y} = \mathbf{x}_{k+1} = \mathbf{x}_k - t\nabla f(\mathbf{x}_k)$ in the quadratic upper bound,

$$\begin{aligned} f(\mathbf{x}_{k+1}) &\leq f(\mathbf{x}_k) - t\|\nabla f(\mathbf{x}_k)\|^2 + \frac{L}{2}t^2\|\nabla f(\mathbf{x}_k)\|^2 \\ &= f(\mathbf{x}_k) - t \left(1 - \frac{Lt}{2}\right) \|\nabla f(\mathbf{x}_k)\|^2 \end{aligned}$$

Note. If $\nabla f(\mathbf{x}_k) \neq 0$ and $0 < t < \frac{2}{L}$, then $f(\mathbf{x}_{k+1}) < f(\mathbf{x}_k)$, so gradient descent with step size $t \in (0, 2/L)$ is indeed a descent method.

Note. We can lower bound the decrease in function value in each step. In particular, for $0 < t \leq \frac{1}{L}$,

$$f(\mathbf{x}_k) - f(\mathbf{x}_{k+1}) \geq \frac{t}{2} \|\nabla f(\mathbf{x}_k)\|^2$$

Convergence analysis

Theorem. If f is convex and L -smooth, and $\mathbf{x}^*$ is a minimum of f , then for step size $t \in (0, \frac{1}{L}]$, the sequence $\{\mathbf{x}_k\}$ produced by the gradient descent algorithm satisfies

$$f(\mathbf{x}_k) - f(\mathbf{x}^*) \leq \frac{\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2tk}$$

Notes.

- $f(\mathbf{x}_k) \downarrow f^*$ as $k \rightarrow \infty$.
- Any limiting point of $\mathbf{x}_k$ is an optimal solution.
- The rate of convergence is $O(1/k)$, i.e. # of iterations to guarantee $f(\mathbf{x}_k) - f(\mathbf{x}^*) \leq \epsilon$ is $O(1/\epsilon)$. For $\epsilon = 10^{-p}$, $k = O(10^p)$, exponential in the number of significant digits!
- Faster convergence with larger t ; best $t = \frac{1}{L}$, but L is unknown.
- Good initial guess helps.

Gradient flow

Consider a **gradient flow**, a continuous version of gradient descent,

$$\frac{d}{ds}\tilde{\mathbf{x}}_s = -\nabla f(\tilde{\mathbf{x}}_s)$$

$f(\tilde{\mathbf{x}}_s)$ is nonincreasing,

$$\frac{d}{ds}f(\tilde{\mathbf{x}}_s) = \langle \nabla f(\tilde{\mathbf{x}}_s), \frac{d}{ds}\tilde{\mathbf{x}}_s \rangle = -\|\nabla f(\tilde{\mathbf{x}}_s)\|^2 \leq 0$$

If f is convex and $\mathbf{x}^*$ is a minimum of f ,

$$\frac{d}{ds}\|\tilde{\mathbf{x}}_s - \mathbf{x}^*\|^2 = 2\langle \tilde{\mathbf{x}}_s - \mathbf{x}^*, \frac{d}{ds}\tilde{\mathbf{x}}_s \rangle = 2\langle \nabla f(\tilde{\mathbf{x}}_s), \mathbf{x}^* - \tilde{\mathbf{x}}_s \rangle \leq 2[f(\mathbf{x}^*) - f(\tilde{\mathbf{x}}_s)]$$

Integrating from 0 to T and using $f(\tilde{\mathbf{x}}_s) \geq f(\tilde{\mathbf{x}}_T)$,

$$\|\tilde{\mathbf{x}}_T - \mathbf{x}^*\|^2 - \|\tilde{\mathbf{x}}_0 - \mathbf{x}^*\|^2 \leq \int_0^T 2[f(\mathbf{x}^*) - f(\tilde{\mathbf{x}}_s)]ds \leq 2T[f(\mathbf{x}^*) - f(\tilde{\mathbf{x}}_T)]$$

so

$$f(\tilde{\mathbf{x}}_T) - f(\mathbf{x}^*) \leq \frac{\|\tilde{\mathbf{x}}_0 - \mathbf{x}^*\|^2}{2T}$$

Proof

Discretizing the gradient flow by Euler's method with step size t ,

$$\frac{\tilde{\mathbf{x}}_{(k+1)t} - \tilde{\mathbf{x}}_{kt}}{t} \approx -\nabla f(\tilde{\mathbf{x}}_{kt})$$

1. By the basic gradient step $\mathbf{x}_{k+1} = \mathbf{x}_k - t\nabla f(\mathbf{x}_k)$,

$$\begin{aligned}\|\mathbf{x}_{k+1} - \mathbf{x}^*\|^2 - \|\mathbf{x}_k - \mathbf{x}^*\|^2 &= \|\mathbf{x}_k - t\nabla f(\mathbf{x}_k) - \mathbf{x}^*\|^2 - \|\mathbf{x}_k - \mathbf{x}^*\|^2 \\ &= 2t\nabla f(\mathbf{x}_k)^T(\mathbf{x}^* - \mathbf{x}_k) + t^2\|\nabla f(\mathbf{x}_k)\|^2\end{aligned}$$

2. By the first-order condition for convexity, the first term is upper bounded by

$$2t\nabla f(\mathbf{x}_k)^T(\mathbf{x}^* - \mathbf{x}_k) \leq 2t[f(\mathbf{x}^*) - f(\mathbf{x}_k)]$$

3. By L -smoothness (the last inequality on slide 15), the second (extra!) term is upper bounded by

$$t^2\|\nabla f(\mathbf{x}_k)\|^2 \leq 2t[f(\mathbf{x}_k) - f(\mathbf{x}_{k+1})]$$

Proof (cont'd)

4. Plugging 2 and 3 into 1,

$$\|\mathbf{x}_{k+1} - \mathbf{x}^*\|^2 - \|\mathbf{x}_k - \mathbf{x}^*\|^2 \leq 2t[f(\mathbf{x}^*) - f(\mathbf{x}_{k+1})]$$

5. Summing over k from 0 to $N - 1$,

$$\|\mathbf{x}_N - \mathbf{x}^*\|^2 - \|\mathbf{x}_0 - \mathbf{x}^*\|^2 \leq 2t \sum_{k=0}^{N-1} [f(\mathbf{x}^*) - f(\mathbf{x}_{k+1})]$$

6. By the descent property $f(\mathbf{x}_{k+1}) \leq f(\mathbf{x}_k)$,

$$\|\mathbf{x}_N - \mathbf{x}^*\|^2 - \|\mathbf{x}_0 - \mathbf{x}^*\|^2 \leq 2tN[f(\mathbf{x}^*) - f(\mathbf{x}_N)]$$

so

$$f(\mathbf{x}_N) - f(\mathbf{x}^*) \leq \frac{\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2tN}$$